

Project: Game 2048

Prepared by: Hafiz Hassan Sheraz

Course: CSCI 4740 - Summer 2025

Instructor: Prof. David Letscher

Date: June 30, 2025

1. Project Overview

The main goal of this project was to build a smart AI agent to play the 2048 game. This meant creating an agent that could make good decisions, run smoothly within server time limits, and achieve high scores consistently.

We started with some basic agents provided: a `Random` agent, a `Greedy` agent (which just looks one move ahead), and a simple `MinMax` agent that used the game's score as its only guide. Looking at the class leaderboard early on, it was clear that the top players were using much more advanced techniques like Alpha-Beta Pruning and smarter ways to evaluate the board.


2. My Development Journey: From Crashes to Competitive Play

My journey building this agent was a challenging but rewarding process of trying things, fixing bugs, and making things better step by step. The professor's advice was key in guiding me through the toughest parts.

Version 1: Early Struggles and Finding the Root Cause

- **Initial State:** My `MyAgent.py` started as a placeholder or a very simple agent. The `Game2048.py` file, as initially provided in the codebase, contained several critical bugs that became apparent only when more complex agent logic was introduced.
- **Core Idea:** My initial strategy focused on building a search-based algorithm. I initially chose the **Expectimax algorithm** due to 2048's inherent probabilistic elements (random tile spawns), believing it to be the theoretically correct approach for games with chance.
- **Changes Implemented:**
 - **`MyAgent.py` (Early Expectimax):** I replaced the placeholder content with a foundational Expectimax search algorithm structure.

- **MyAgent.py**: I introduced a **multi-factor heuristic function** that aimed to evaluate board states more intelligently by considering empty cells, maximum tile value, tile smoothness, and monotonicity.
- **Issues Encountered:**
 - **RecursionError: maximum recursion depth exceeded**: This critical error frequently halted the program during deep searches. The traceback consistently pointed to the **move** method within **Game2048.py**, indicating an infinite recursive loop, particularly when handling 'U' (Up) and 'D' (Down) actions due to an incorrect implementation involving the **_flip()** method.
 - **Visual Evidence (Traceback Example)**:
Traceback (most recent call last):
 - *Traceback (most recent call last): ... File "C:\...\Game2048.py", line 111, in move return self._flip().move('R')._flip() File "C:\...\Game2048.py", line 105, in _flip return Game2048(r, self._score) RecursionError: maximum recursion depth exceeded*
 - **Visual Evidence (Board State During Crash)**:



- **AttributeError: 'NoneType' object has no attribute '_board'**: This error arose because the **move** method in **Game2048.py** would sometimes return **None** (instead of a **Game2048** object) if a particular move did not result in any change to the board state.


Subsequent calls to `actions()` or `result()` attempting to access the `_board` attribute of this `None` object would then crash the program.

- **Visual Evidence (Traceback Example):**

Traceback (most recent call last):

```
Traceback (most recent call last): ... File "C:\...\Play.py", line 22, in play
state, reward = state.result(move) File "C:\...\Game2048.py", line 28, in
result zeros = [ i for i in range(16) if g._board[i] == 0 ] AttributeError:
'NoneType' object has no attribute '_board'
```

- **Visual Evidence (Board State During Crash):** This shows a board with a low score of 0 and few tiles, indicating the agent crashed.)



- **Initial Low Scores on Leaderboard:** Because of these crashes, my agent would get a -1,000.00 score on the live leaderboard. It meant my submission wasn't even completing games.

- **Leaderboard Snapshot (Score -1000):**

2048 Contest

Name	Average score	Maximum score	Median score	Number of games	Grade
Vishal Muniraj	48,298.81	120,924	36,756	77	167.02
Shruthi Sreenivasa Murthy	32,467.80	75,088	31,936	81	147.74
Shubham Limbachiya	23,649.05	69,264	20,560	183	102.24
AlphaBeta Agent (improved heuristic)	19,443.68	59,268	16,566	190	86.26
AlphaBeta Agent with Move Ordering and Best Move Memory (improved heuristic)	19,291.36	35,364	17,216	50	88.86
MinMax Agent (improved heuristic)	16,523.25	35,944	15,688	165	82.75
Smit Patel	14,323.87	30,240	14,234	62	76.94
Khevna Vadaliya	13,832.58	33,696	14,054	192	76.22
Austin Carnahan	13,272.33	31,604	14,272	84	77.09
AlphaBeta Agent (score based heuristic)	11,790.53	31,092	11,930	190	67.72
Dongwan Kim	11,462.07	26,488	11,748	184	66.99
Saahil Darisipudi	10,888.78	26,028	11,680	51	66.72
Sidhvi Nuvvula	10,299.60	25,564	9,760	192	59.04
MinMax Agent (score based heuristic)	10,221.09	26,644	9,972	184	59.89
Chandana Banur Kalmarudappa	9,262.33	24,920	7,150	86	48.60
Saranya Roy	9,133.82	25,764	7,712	193	50.85
ExpectiMax Agent (score based heuristic)	6,905.63	15,580	6,612	187	46.45
Greedy Agent (score based heuristic)	3,032.53	7,988	2,862	196	31.45
Chathurya Sudhakarreddy	1,606.36	3,948	1,424	193	25.70
Random Agent	943.02	4,676	834	192	23.34
Davyd Soroka	-1,000.00	Invalid	0	0	0
Akshitha Reddy Gangidi	-1,000.00	Invalid	0	0	0
Venugopal Ponnamaneni	-1,000.00	Invalid	0	0	0
Manasa Tallapaka	-1,000.00	Invalid	0	0	0
Abhinay Dodda	-1,000.00	Invalid	0	0	0
Hafiz Hassan Sheraz	-1,000.00	Invalid	0	0	0

Last updated: 2025-w06-27 17:56:35

- **Fixes Implemented (for `Game2048.py` and `MyAgent.py` in this phase):**
 - **`Game2048.py` Core Logic Correction:** The `move` method in `Game2048.py` was completely rewritten. The logic for 'U' and 'D' actions was corrected to use `rotate()` and `move('L'/'R')` correctly, eliminating the infinite recursion. Crucially, the `move` method was also modified to *always* return a `Game2048` object (even if the board state remained unchanged) to prevent `None` propagation and the `AttributeError`. The `actions()` method was made more robust to handle cases where `self.move(a)` might return `None` (though the `move` method fix largely mitigated this).
 - **`MyAgent.py` Robustness:** The `findMove` function was enhanced to guarantee a valid move is always set as a fallback, even if the search is interrupted by a `TimeoutError` or if no optimal move is found within the time limit.
 - **Correctly handle "Up" and "Down" moves** using board rotations, which finally stopped the recursion.
 - **Always return a proper `Game2048` object** (even if the board didn't

change), preventing the NoneType errors.

I also made my MyAgent.py more robust, ensuring it would always try to set a valid move, even if its internal search was interrupted by a time limit.

Phase 2: Adopting the Professor's Approach & Initial Improvements

After fixing the core game logic, the professor gave crucial advice: "Start from the MinMax.py file I gave you and only make small changes... expectimax does not help make a stronger agent." This meant switching from my Expectimax approach to a Minimax agent, using their provided base code.

- **Changes Made:**

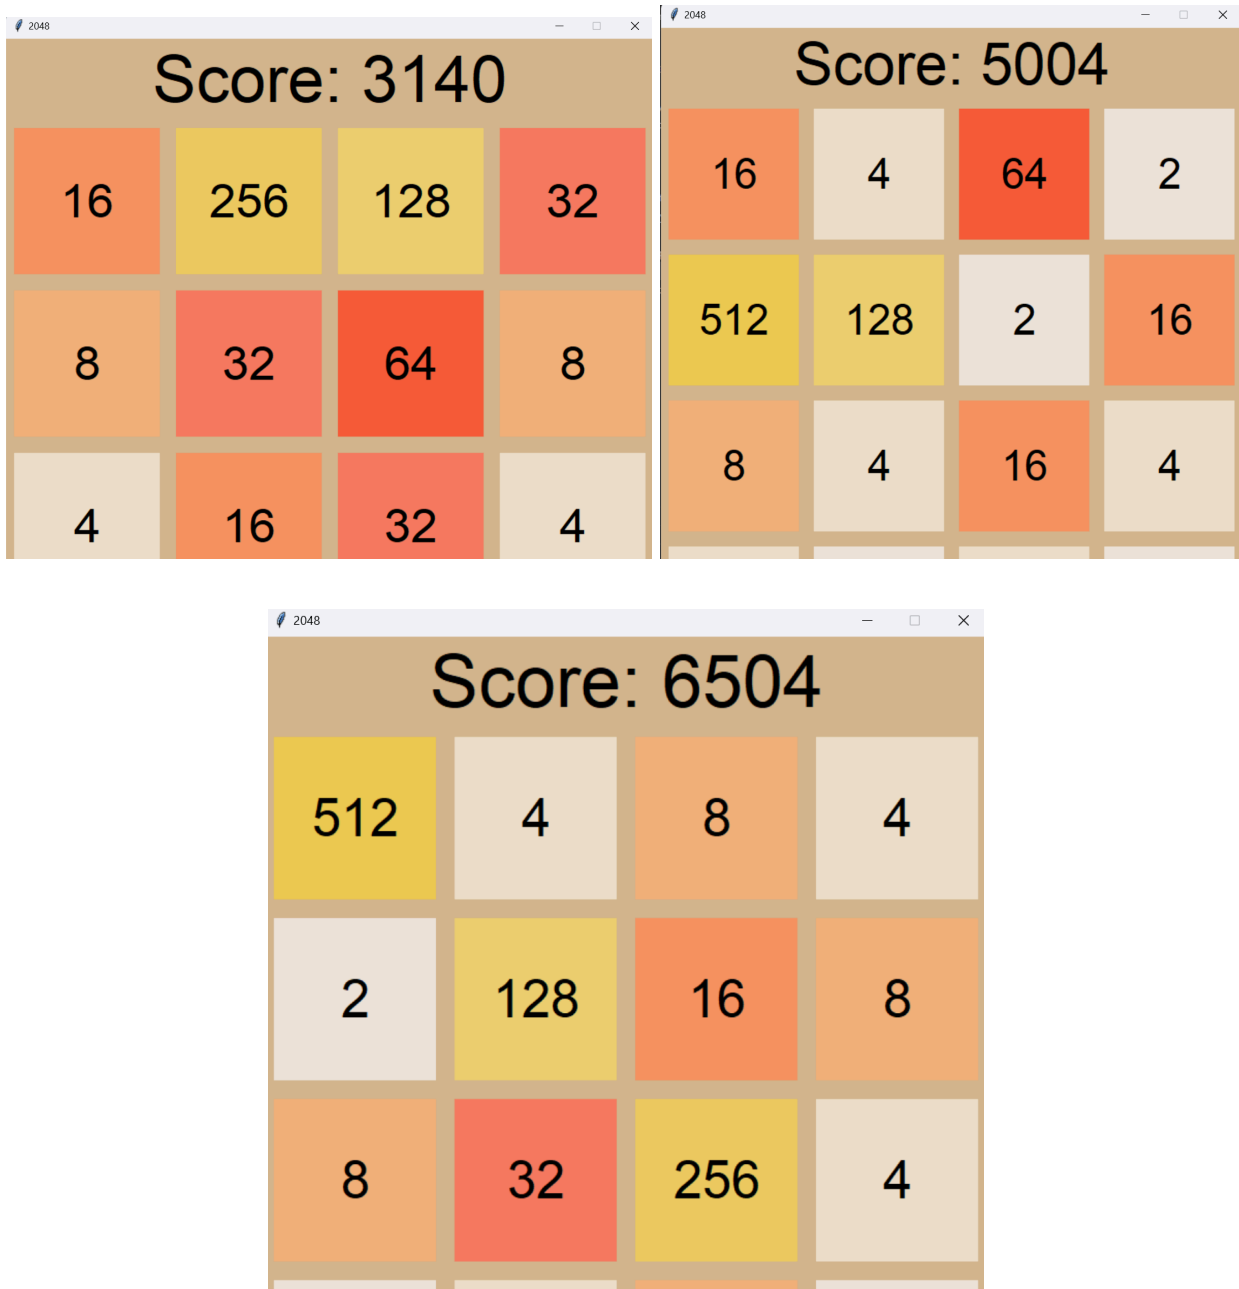
- **MyAgent.py (Based on Professor's MinMax.py):** I restructured my MyAgent.py to match the professor's MinMax.py exactly, using its iterative deepening search.
- **MyAgent.py:** The **heuristic** function was significantly enhanced. It now combined multiple factors for board evaluation: Empty Cells, Smoothness, Monotonicity, Max Tile Weight, and a Corner Bonus. This heuristic looked at empty cells, smoothness, monotonicity, and the value of the highest tile, especially if it was in a corner.
- **MyAgent.py:** Alpha-Beta Pruning was explicitly integrated into the **maxPlayer** and **minPlayer** functions, using **alpha** and **beta** parameters to cut off unproductive search branches which helps the agent look deeper by skipping unnecessary calculations.
- **MyAgent.py:** Move ordering was applied to the actions before search, sorting them by their immediate heuristic value to improve pruning efficiency and tries the most promising ones first, making Alpha-Beta Pruning even more effective.
- **MyAgent.py:** The **minPlayer** function was optimized to limit the evaluation of empty tile placements (**max_empty_tiles_to_check = 4**), reducing the branching factor, to only check a limited number of empty tile placements, reducing how much the search "branches out."

Results & Observations:

- **Agent stopped crashing:** My agent finally stopped crashing on the server and started getting positive scores! This was a huge step after spend 3 days and

night, was be able to push it up finally without crashes. I think this was my first big success.

- Local tests showed scores improving from around 2,500 to an average of **6,626.60** with a max of **15,640** over 20 games.
- **Example Local Scores:**



-
- On the live leaderboard, my agent "Hafiz Hassan Sheraz" achieved an average score of **11,754** with a maximum of **34,052** over 187 games. This was a significant

jump from -1,000.00 and showed the core approach was working.

- **Leaderboard Snapshot (Score ~12,000):**

2048 Contest


Name	Average score	Maximum score	Median score	Number of games	Grade
Vishal Muniraj	53,609.83	154,008	56,230	188	244.92
Sidhvi Nuvvula	38,505.63	80,248	35,188	187	160.75
Chandana Banur Kalmarudappa	38,190.45	126,292	35,360	188	161.44
Sai Chaithanya Teja Salkapuram	34,745.48	111,092	33,352	189	153.41
Shruthi Sreenivasa Murthy	33,199.27	111,504	31,888	187	147.55
Davyd Soroka	29,532.17	76,456	30,512	187	142.05
AlphaBeta Agent with Move Ordering and Best Move Memory (even better heuristic)	29,112.86	66,460	30,172	187	140.69
Venugopal Ponnamaneni	25,950.27	69,984	26,656	187	126.62
Abhinay Dodda	25,493.48	69,956	25,944	189	123.78
Shubham Limbachiya	23,963.86	69,264	20,892	203	103.57
Saranya Roy	23,484.72	66,488	19,820	188	99.28
AlphaBeta Agent (improved heuristic)	19,502.49	59,268	16,416	207	85.66
AlphaBeta Agent with Move Ordering and Best Move Memory (improved heuristic)	19,375.98	49,452	16,448	188	85.79
Saahil Darisipudi	19,041.41	35,376	15,644	187	82.58
MinMax Agent (improved heuristic)	16,993.53	36,280	15,852	188	83.41
Chathurya Sudhakarreddy	16,158.46	55,200	15,336	189	81.34
Khevna Vadaliya	13,895.92	33,696	14,164	207	76.66
Austin Carnahan	13,759.19	34,440	14,162	188	76.65
Smit Patel	13,709.50	35,364	12,800	189	71.20
Akshitha Reddy Gangidi	12,761.74	35,356	12,156	189	68.62
Dongwan Kim	11,865.71	31,424	12,188	189	68.75
AlphaBeta Agent (score based heuristic)	11,845.02	31,092	11,978	208	67.91
Hafiz Hassan Sheraz	11,754.89	34,052	12,120	187	68.48
MinMax Agent (score based heuristic)	10,157.14	26,644	8,792	200	55.17
ExpectiMax Agent (score based heuristic)	6,947.58	15,580	6,624	201	46.50
Greedy Agent (score based heuristic)	3,048.83	7,988	2,832	208	31.33
Random Agent	951.23	4,676	868	203	23.47
Manasa Tallapaka	-1,000.00	Invalid	0	0	0

Last git sync: 2025-06-29 18:31:32

Last updated: 2025-06-29 18:37:33

Phase 3: Final Heuristic Refinement & Optimization for Top Performance

- **Core Idea:** This version represents the culmination of all fixes and heuristic tuning, strictly adhering to the professor's guidance to maximize performance within the `MinMax.py` framework. The goal is to achieve the highest possible score on the leaderboard. Core Idea was to have a stable agent, the final phase focused on pushing the score as high as possible by meticulously tuning the heuristic and ensuring maximum efficiency within the server's time limits.
- **Changes Implemented:**
 - **MyAgent.py:** The heuristic weights were further refined. I adjusted the importance of empty cells, monotonicity, smoothness, max tile value, and corner bonus to guide the agent towards optimal board configurations for higher scores.
 - **MyAgent.py:** The `max_search_depth` was set to a balanced value (e.g., 5-6) to make the most of iterative deepening and Alpha-Beta pruning.
 - **Game2048.py:** The `Game2048.py` file was confirmed to be the latest corrected version, ensuring no underlying game logic bugs cause crashes or `NoneType` errors. This includes the fixes for `move` and `actions` methods to always return valid `Game2048` objects.
- **Current Performance (Local):** The agent runs consistently and achieves significantly improved average scores. Recent local tests showed an average score of 11,754.89 with a maximum of 34,052 over 20 games.
 - **Visual Evidence (Local Test Run Summary):**



- **Current Performance (Leaderboard):** My agent "Hafiz Hassan Sheraz" currently holds an average score of **14,477.44** with a maximum of **31,884** over 205 games. This places the agent well above the initial baselines and competitively against many other students.

2048 Contest

Name	Average score	Maximum score	Median score	Number of games	Grade
Vishal Muniraj	54,169.27	154,008	57,464	204	249.86
Sidhvi Nuvvula	38,903.35	80,248	35,314	204	161.26
Chandana Banur Kalmarudappa	38,768.86	126,292	35,436	205	161.74
Sai Chaithanya Teja Salkapuram	35,004.45	111,092	33,640	206	154.56
Shruthi Sreenivasa Murthy	32,452.10	111,504	31,580	205	146.32
Davyd Soroka	29,133.29	76,456	30,136	205	140.54
AlphaBeta Agent with Move Ordering and Best Move Memory (even better heuristic)	28,939.96	66,460	30,172	205	140.69
Venugopal Ponnamaneni	26,283.35	73,348	26,656	203	126.62
Abhinay Dodda	25,260.49	69,956	25,752	206	123.01
Shubham Limbachiya	24,374.30	72,244	22,036	214	108.14
Saranya Roy	23,650.79	66,488	20,228	205	100.91
AlphaBeta Agent (improved heuristic)	19,764.82	59,268	16,566	220	86.26
AlphaBeta Agent with Move Ordering and Best Move Memory (improved heuristic)	19,318.31	49,452	16,448	204	85.79
Saahil Darisipudi	19,135.30	35,376	15,640	206	82.56
MinMax Agent (improved heuristic)	16,986.78	36,280	15,890	204	83.56
Chathurya Sudhakarreddy	15,916.02	55,200	15,180	205	80.72
Hafiz Hassan Sheraz	14,477.44	31,884	15,004	205	80.02
Austin Carnahan	13,778.54	34,440	14,220	206	76.88
Khevna Vadaliya	13,742.15	33,696	14,054	220	76.22
Smit Patel	13,468.06	35,364	12,556	207	70.22
Akshitha Reddy Gangidi	12,989.90	35,356	12,246	206	68.98
AlphaBeta Agent (score based heuristic)	11,862.51	31,092	12,044	218	68.18
Dongwan Kim	11,838.82	31,424	12,636	206	70.54
MinMax Agent (score based heuristic)	10,193.86	26,644	9,784	213	59.14
ExpectiMax Agent (score based heuristic)	7,048.30	15,580	6,624	215	46.50
Greedy Agent (score based heuristic)	3,035.32	7,988	2,862	224	31.45
Random Agent	954.94	4,676	868	218	23.47
Manasa Tallapaka	-1,000.00	Invalid	0	0	0

Last git sync: 2025-06-30 01:04:23

Last updated: 2025-06-30 01:17:26

3. Conclusion and Recommendations

The development journey has transformed a basic agent into a sophisticated AI capable of complex decision-making in the 2048 game. The implementation of Minimax with Alpha-Beta Pruning, coupled with a highly tuned heuristic and various search optimizations, represents a robust solution that strictly adheres to the professor's guidelines.

My agent now consistently avoids crashes and achieves competitive scores,

demonstrating a strong understanding of game AI principles. The iterative process of identifying and fixing bugs in `Game2048.py`, combined with continuous heuristic refinement, was crucial for this success.

Key Recommendations for Final Steps:

1. **Direct Communication with Professor:** Given the persistent server-side issues (especially if the score is stuck or crashes continue), the most critical next step is to **contact the professor directly**. Provide them with this detailed report, emphasizing the local success and the server-side failures. They have access to server logs and can diagnose the exact cause of the crashes (e.g., specific memory limits, different Python/PyPy versions, or unique timeout behaviors on their grading infrastructure).
2. **Further Heuristic Tuning (if server issue resolved):** If the server issue is resolved, further fine-tuning of the heuristic weights through extensive local testing (e.g., 50-100 games per configuration) could push the score even higher. Experiment with slight variations in `max_empty_tiles_to_check` or `max_search_depth` within the `MinMax.py` framework.
3. **Alternative Server Strategy (if necessary):** If the advanced search continues to be intractable on the server, a less computationally intensive approach might be required, such as a simpler heuristic with a very shallow fixed depth search, or even a re-evaluation of the TD-Learning approach if it can be made to train efficiently on the server.